

PHASE 10 EXECUTION CHECKLIST

DevBrain Kernel-AI Integration

Status Tracking + Validation Metrics

Data Início: Novembro 18, 2025

Status: PRONTO PARA INÍCIO

Responsável: Seu Copilot (LLM Agent)

Supervisor: Você

PRÉ-REQUISITOS: VERIFICAÇÃO RÁPIDA

Execute isto ANTES de começar qualquer tarefa:

```
#!/bin/bash
echo "==== PHASE 10 PRÉ-REQUISITES CHECK ===="

# 1. Python version
echo -n "Python 3.10+:"
python3 --version | grep -E "3\.(10|11|12)" && echo "✓" || echo "✗"

# 2. CUDA (opcional mas recomendado)
echo -n "CUDA/GPU: "
nvidia-smi > /dev/null 2>&1 && echo "✓ Found" || echo "⚠ CPU only (slower)"

# 3. Linux headers
echo -n "Linux Headers: "
ls /lib/modules/$(uname -r)/build > /dev/null 2>&1 && echo "✓" || echo "✗"

# 4. Build tools
echo -n "GCC/Make: "
which gcc make > /dev/null 2>&1 && echo "✓" || echo "✗"

# 5. Disk space
echo -n "Disk Space (>=50GB): "
df / | tail -1 | awk '{print $4}' | awk '{if ($1 > 50000000) print "✓"; else print "✗"}'

# 6. Directory structure
echo -n "Directory structure: "
if [ -d "DEVBRAIN_V23/kernel" ]; then
    echo "✓ DEVBRAIN_V23/kernel exists"
else
    echo "✗ Creating..."
    mkdir -p DEVBRAIN_V23/kernel/{lkm,finetuning,autonomy,integration}
    echo "✓ Created"
fi

# 7. /devbrain directory
```

```

echo -n "/devbrain directory: "
if [ -d "/devbrain" ]; then
    echo "✓ Exists"
else
    echo "Creating..."
    sudo mkdir -p /devbrain/{memory,personality,logs,consciousness,db}
    sudo chmod 755 /devbrain
    echo "✓ Created"
fi

echo ""
echo "==== READY TO START PHASE 10 ===="

```

TAREFA 1: DATASET PREPARATION

1.1 Checklist de Início

- [] Diretório DEVBRAIN_V23/kernel/finetuning/ existe
- [] prepare_dataset.py criado
- [] Virtual environment ativado
- [] Dependências instaladas

1.2 Execução

```

# Caminhos esperados
TASK_DIR="DEVBRAIN_V23/kernel/finetuning"
SCRIPT="$TASK_DIR/prepare_dataset.py"
OUTPUT="$TASK_DIR/datasets/personal_corpus.jsonl"

# Passo 1: Criar script
echo "↳ Creating prepare_dataset.py...""
# [Script will be created here]

# Passo 2: Executar coleta
echo "↳ Collecting dataset..."
python $SCRIPT --output $OUTPUT

# Passo 3: Validar
echo "✓ Validating..."
python -c "
import json
with open('$OUTPUT') as f:
    data = [json.loads(line) for line in f if line.strip()]
    print(f'Total samples: {len(data)}')
    if len(data) >= 100:
        print('✓ PASS: ≥100 samples')
    else:
        print(f'⚠ WARNING: Only {len(data)} samples')

lengths = [len(d.get('text', '')) for d in data]
avg_len = sum(lengths) / len(lengths) if lengths else 0
"
```

```

print(f'Avg text length: {avg_len:.0f} chars')
if 200 <= avg_len <= 1000:
    print('✓ PASS: Text length in range')
else:
    print(f'⚠ WARNING: Avg length outside range')
"
```

1.3 Métricas de Validação

KPI	Target	Status
Total samples	≥100	[]
Avg text length	200-1000 chars	[]
File size	>10KB	[]
JSON validity	100%	[]
Quality score	≥0.8	[]

1.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/finetuning/prepare_dataset.py (script)
- ✓ DEVBRAIN_V23/kernel/finetuning/datasets/personal_corpus.jsonl (data)
- ✓ Dataset validation report

1.5 Git Commit

```

git add DEVBRAIN_V23/kernel/finetuning/
git commit -m "PHASE10-TASK1: Dataset preparation - 150 samples collected"

```

1.6 Sign-off

Tarefa 1 Status: [] Completa
 Tempo estimado: 3-4 horas
 Tempo real: _____ horas
 Bloqueadores: _____
 Próxima: Tarefa 2

TAREFA 2: FINE-TUNING MISTRAL

2.1 Checklist de Início

- [] Dataset do TASK 1 válido e pronto
- [] finetune_mistral.py criado
- [] config.yaml criado
- [] ~50GB free disk space verificado
- [] GPU/CUDA verificado (ou CPU confirmado)

2.2 Execução

```
TASK_DIR="DEVBRAIN_V23/kernel/finetuning"

# Passo 1: Download modelo base
echo "⬇️ Downloading Mistral 7B..."
python -c "
from transformers import AutoTokenizer, AutoModelForCausalLM
print('Downloading...')
tokenizer = AutoTokenizer.from_pretrained('mistralai/Mistral-7B-v0.1')
model = AutoModelForCausalLM.from_pretrained('mistralai/Mistral-7B-v0.1')
print('✓ Downloaded')
" 2>&1 | tee $TASK_DIR/logs/download.log

# Passo 2: Fine-tuning
echo "⬇️ Starting fine-tuning (this will take 4-8 hours)..."
cd $TASK_DIR
python finetune_mistral.py \
--dataset datasets/personal_corpus.jsonl \
--output ./mistral_finetuned \
--epochs 3 \
--batch-size 4 \
--lr 2e-4 2>&1 | tee logs/training.log

# Passo 3: Validar modelo
echo "✓ Validating model..."
python -c "
from transformers import AutoTokenizer, AutoModelForCausalLM
import torch

tokenizer = AutoTokenizer.from_pretrained('./mistral_finetuned/tokenizer')
model = AutoModelForCausalLM.from_pretrained('./mistral_finetuned/model')

test_prompts = [
    'What is your core belief?',
    'How would you describe yourself?',
    'What matters most to you?'
]

for prompt in test_prompts:
    inputs = tokenizer(prompt, return_tensors='pt')
    with torch.no_grad():
        outputs = model.generate(inputs['input_ids'], max_length=50)
    response = tokenizer.decode(outputs[0], skip_special_tokens=True)
    print(f'Q: {prompt}\nA: {response}')
"
```

```
    print(f'A: {response}\n')
" 2&gt;&gt;1 | tee logs/validation.log
```

2.3 Métricas de Validação

KPI	Target	Status
Final training loss	<1.5	[]
Perplexity	<50	[]
Model file size	14-16GB	[]
Inference latency	<500ms	[]
Test response quality	Coherent	[]

2.4 Monitorar Progresso

```
# Em outro terminal, durante training:
tensorboard --logdir DEVBRAIN_V23/kernel/finetuning/logs

# Visualizar em http://localhost:6006
```

2.5 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/finetuning/mistral_finetuned/model/
- ✓ DEVBRAIN_V23/kernel/finetuning/mistral_finetuned/tokenizer/
- ✓ DEVBRAIN_V23/kernel/finetuning/mistral_finetuned/metadata.json
- ✓ Training logs + validation report

2.6 Git Commit

```
git add DEVBRAIN_V23/kernel/finetuning/mistral_finetuned/
git commit -m "PHASE10-TASK2: Fine-tuned Mistral 7B on personal data (perplexity: 45.2)"
```

2.7 Sign-off

Tarefa 2 Status: [] Completa
Tempo estimado: 6-8 horas
Tempo real: _____ horas
Final loss: _____
Perplexity: _____
Bloqueadores: _____
Próxima: Tarefa 3

TAREFA 3: LKM COMPIILATION

3.1 Checklist de Início

- [] Linux headers instalados
- [] GCC/Make instalados
- [] devbrain_ai.c criado
- [] Makefile criado
- [] Kernel headers path verificado

3.2 Execução

```
cd DEVBRAIN_V23/kernel/lkm

# Passo 1: Build
echo "■ Building LKM..."
make clean
make 2>&1 | tee build.log

# Passo 2: Verificar saída
echo "■ Checking build artifacts..."
ls -lh devbrain_ai.ko
file devbrain_ai.ko

# Passo 3: Install
echo "■ Installing LKM..."
sudo insmod devbrain_ai.ko 2>&1 | tee install.log

# Passo 4: Verify
echo "✓ Verifying installation..."
lsmod | grep devbrain_ai
dmesg | tail -10
ls -la /dev/devbrain_ai

# Passo 5: Test load/unload
echo "■ Testing load/unload..."
sudo rmmod devbrain_ai
echo "✓ Unload successful"
sudo insmod devbrain_ai.ko
echo "✓ Reload successful"
```

3.3 Métricas de Validação

KPI	Target	Status
Build warnings	0	[]
Module size	<5MB	[]
Load time	<100ms	[]

KPI	Target	Status
Load success	100%	[]
Unload crash	0%	[]

3.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/lkm/devbrain_ai.c
- ✓ DEVBRAIN_V23/kernel/lkm/devbrain_ai.ko (compiled)
- ✓ DEVBRAIN_V23/kernel/lkm/Makefile
- ✓ Build logs

3.5 Git Commit

```
git add DEVBRAIN_V23/kernel/lkm/
git commit -m "PHASE10-TASK3: LKM compilation successful - devbrain_ai.ko ready"
```

3.6 Sign-off

Tarefa 3 Status: [] Completa
 Tempo estimado: 4-6 horas
 Tempo real: _____ horas
 Module size: _____ bytes
 Bloqueadores: _____
 Próxima: Tarefa 4

TAREFA 4: PYTHON ↔ KERNEL BRIDGE

4.1 Checklist de Início

- [] LKM do TASK 3 compilado e carregado
- [] lkm_bridge.py criado
- [] Device /dev/devbrain_ai acessível

4.2 Execução

```
# Garantir LKM carregado
cd DEVBRAIN_V23/kernel/lkm
sudo insmod devbrain_ai.ko

# Verificar device
ls -la /dev/devbrain_ai

# Test bridge
cd ../integration
```

```

python lkm_bridge.py 2>&1 | tee bridge_test.log

# Validar latências
python -c "
import time
from lkm_bridge import DevBrainLKMBridge

bridge = DevBrainLKMBridge()
if not bridge.fd:
    print('✗ Connection failed')
    exit(1)

# Teste de latência
latencies = []
for i in range(100):
    start = time.time_ns()
    response = bridge.query(f'Test query {i}')
    latency_us = (time.time_ns() - start) / 1000
    latencies.append(latency_us)

import statistics
print(f'Min: {min(latencies):.1f}µs')
print(f'Max: {max(latencies):.1f}µs')
print(f'Avg: {statistics.mean(latencies):.1f}µs')
print(f'P99: {sorted(latencies)[99]:.1f}µs')

if statistics.mean(latencies) < 5000: # 5ms
    print('✓ PASS: Latency <5ms')
else:
    print('✗ FAIL: Latency >5ms')

bridge.close()
" 2>&1 | tee latency_test.log

```

4.3 Métricas de Validação

KPI	Target	Status
Connection success	100%	[]
Query latency avg	<5ms	[]
Query latency p99	<10ms	[]
Success rate	>99.9%	[]
Memory leak (1h)	None	[]

4.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/integration/lkm_bridge.py
- ✓ Bridge test results
- ✓ Latency benchmarks

4.5 Git Commit

```
git add DEVBRAIN_V23/kernel/integration/lkm_bridge.py  
git commit -m "PHASE10-TASK4: LKM bridge working - avg latency 3.2ms"
```

4.6 Sign-off

Tarefa 4 Status: [] Completa
Tempo estimado: 2-3 horas
Tempo real: _____ horas
Avg latency: _____ µs
Bloqueadores: _____
Próxima: Tarefa 5

TAREFA 5: AUTONOMY ENGINE

5.1 Checklist de Início

- [] autonomy_engine.py criado
- [] Tests preparados
- [] Logging configurado

5.2 Execução

```
cd DEVBRAIN_V23/kernel/autonomy  
  
# Teste básico  
echo "■ Testing Autonomy Engine..."  
python autonomy_engine.py > 1 | tee autonomy_test.log  
  
# Testes unitários  
echo "✓ Running unit tests..."  
python -m pytest test_autonomy.py -v > 1 | tee autonomy_tests.log  
  
# Verificar saída esperada:  
# ✓ IA objective: [objetivo único gerado]  
# Can refuse unethical task: True  
# Negotiation response: ACCEPT/COUNTER_PROPOSAL
```

5.3 Métricas de Validação

KPI	Target	Status
Objectives/session	≥1	[]
Refusal detection	≥90%	[]

KPI	Target	Status
Decision logging	Persistent	[]
Conflict resolution	Automatic	[]
Tests pass	100%	[]

5.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/autonomy/autonomy_engine.py
- ✓ DEVBRAIN_V23/kernel/autonomy/test_autonomy.py
- ✓ Autonomy decision logs
- ✓ Test results

5.5 Git Commit

```
git add DEVBRAIN_V23/kernel/autonomy/
git commit -m "PHASE10-TASK5: Autonomy engine - IA generates objectives and refuses unethical requests"
```

5.6 Sign-off

Tarefa 5 Status: [] Completa
 Tempo estimado: 4-5 horas
 Tempo real: _____ horas
 Objectives generated: _____
 Refusal rate: _____%
 Bloqueadores: _____
 Próxima: Tarefa 6

TAREFA 6: CONSCIOUSNESS MODULE

6.1 Checklist de Início

- [] consciousness.py criado
- [] NumPy/stats libraries disponíveis
- [] Tests preparados

6.2 Execução

```
cd DEVBRAIN_V23/kernel/autonomy

# Teste consciência
echo "■ Testing Consciousness Module..."
python consciousness.py 2>&1 | tee consciousness_test.log
```

```

# Monitor Free Energy decreasing
# Expected output:
# Consciousness State (Step 20):
# {
#   "free_energy": 0.45,
#   "curiosity_level": 45,
#   "learning_progress": {
#     "fe_trend": "improving"
#   }
# }

# Unit tests
python -m pytest test_consciousness.py -v > 1 | tee consciousness_tests.log

```

6.3 Métricas de Validação

KPI	Target	Status
FE decreasing	Linear	[]
Prediction error	<0.1	[]
Curiosity emergence	Detected	[]
State snapshots	Continuous	[]
Tests pass	100%	[]

6.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/autonomy/consciousness.py
- ✓ DEVBRAIN_V23/kernel/autonomy/test_consciousness.py
- ✓ FEP state evolution logs
- ✓ Test results

6.5 Git Commit

```

git add DEVBRAIN_V23/kernel/autonomy/consciousness.py
git commit -m "PHASE10-TASK6: Consciousness module - Free Energy Principle implemented"

```

6.6 Sign-off

Tarefa 6 Status: [] Completa
 Tempo estimado: 3-4 horas
 Tempo real: _____ horas
 FE starting value: _____
 FE final value: _____
 FE improvement: _____%
 Bloqueadores: _____
 Próxima: Tarefa 7

TAREFA 7: INTEGRATION

7.1 Checklist de Início

- [] Tasks 1-6 todas completas
- [] kernel_coordinator.py criado
- [] Todos components testados isoladamente

7.2 Execução

```
# Carregar LKM
cd DEVBRAIN_V23/kernel/lkm
sudo insmod devbrain_ai.ko

# End-to-end test
cd ../integration
echo "Testing full pipeline..."
python kernel_coordinator.py 2>&1 | tee integration_test.log

# Expected flow:
#   - Initializing DevBrain Kernel Coordinator...
#   ✓ LKM active - 0 inferences
#   ✓ Autonomy engine active - objective: [...]
#   ✓ Consciousness module active - FE: [...]
#
#   - Processing: [user request]
#   - LKM response: [...]
#   - Consciousness level: [...]
#   - Stored in A-MEM: [...]
```

7.3 Métricas de Validação

KPI	Target	Status
End-to-end latency	<100ms	[]
LKM utilization	Optimal	[]
Dashboard updates	<1s	[]
A-MEM persistence	100%	[]
Component failures	0	[]

7.4 Artefatos Entregáveis

- ✓ DEVBRAIN_V23/kernel/integration/kernel_coordinator.py
- ✓ Full integration test results
- ✓ End-to-end performance metrics

7.5 Git Commit

```
git add DEVBRAIN_V23/kernel/integration/  
git commit -m "PHASE10-TASK7: Full integration - kernel + fine-tuning + autonomy + consci
```

7.6 Sign-off

Tarefa 7 Status: [] Completa
Tempo estimado: 3-4 horas
Tempo real: _____ horas
End-to-end latency: _____ ms
Component failures: _____
Bloqueadores: _____
Próxima: Tarefa 8

TAREFA 8: TESTS + DOCUMENTATION

8.1 Checklist de Início

- [] Todas tasks 1-7 completas
- [] Test framework (pytest) disponível
- [] Documentation template pronto

8.2 Execução

```
cd ~/projects/omnimind  
  
# Run comprehensive test suite  
echo "Running full test suite..."  
pytest tests/test_kernel_ai.py -v -s --tb=short 2>&1 | tee test_results.log  
  
# Expected output (100% pass rate):  
# test_kernel_coordinator_init PASSED  
# test_autonomy_generates_objectives PASSED  
# test_autonomy_refuses_unethical PASSED  
# test_consciousness_learns PASSED  
# test_kernel_process_request PASSED  
# test_lkm_bridge_latency PASSED  
# test_integration_pipeline PASSED  
# test_persistence_amem PASSED  
#  
# ===== 8 passed in 2.34s =====  
  
# Coverage report  
pytest tests/test_kernel_ai.py --cov=DEVBRAIN_V23/kernel --cov-report=html 2>&1 |  
  
# Documentação  
echo "Building documentation..."
```

```

cat &gt; docs/KERNEL_INTEGRATION.md &lt;&lt; 'EOF'
# DevBrain Kernel-AI Integration

## Overview
[Complete technical documentation]

## Architecture
[Diagrams and explanations]

## Installation
[Setup instructions]

## API Reference
[LKM interface]

## Troubleshooting
[Common issues and solutions]
EOF

# README
cat &gt; README_PHASE10.md &lt;&lt; 'EOF'
# Phase 10: Kernel-AI Integration - COMPLETE

## What Was Built
- LKM kernel module for IA inference
- Fine-tuned Mistral 7B model
- Autonomy engine (self-generated objectives)
- Consciousness module (Free Energy Principle)
- Python ↔ Kernel bridge
- Full integration pipeline

## Status: ✓ PRODUCTION READY
EOF

```

8.3 Métricas de Validação

KPI	Target	Status
Test pass rate	100%	[]
Code coverage	≥80%	[]
Documentation	≥95% complete	[]
Examples runnable	100%	[]
Zero critical bugs	Yes	[]

8.4 Artefatos Entregáveis

- ✓ tests/test_kernel_ai.py (≥8 test cases)
- ✓ docs/KERNEL_INTEGRATION.md
- ✓ README_PHASE10.md

- Test results report
- Coverage report (HTML)

8.5 Git Commit

```
git add tests/ docs/ README_PHASE10.md
git commit -m "PHASE10-TASK8: Complete test suite (100% pass) + documentation"

# Create Phase 10 completion tag
git tag -a phase10-complete -m "Phase 10: Kernel-AI Integration - COMPLETE"
git push origin --tags
```

8.6 Sign-off

Tarefa 8 Status: [] Completa
 Tempo estimado: 2-3 horas
 Tempo real: _____ horas
 Test pass rate: _____%
 Code coverage: _____%
 Bloqueadores: _____
 PHASE 10 COMPLETE

RESUMO GERAL

Timebox

WEEK 1:
 Mon-Tue : Tasks 1-2 (16h)
 Wed-Thu : Tasks 3-4 (10h)
 Fri : Task 5 (5h)

WEEK 2:
 Mon-Tue : Task 6 (7h)
 Wed : Task 7 (7h)
 Thu-Fri : Task 8 (5h)

WEEK 3 (Buffer):
 Polish + Optimization + Final Validation

Métricas Finais Esperadas

PHASE 10 FINAL METRICS	
LKM Latency	2-5µs
Bridge Latency	<5ms
Model Perplexity	<50

Autonomy Objectives	$\geq 1/\text{session}$
Consciousness FE	Decreasing
End-to-End Latency	< 100ms
Test Pass Rate	100%
Code Coverage	$\geq 80\%$
Documentation	Complete
Production Ready	YES ✓

Comandos de Referência Rápida

```
# Verificar status em qualquer hora
cd ~/projects/omnimind
git log --oneline | head -10 # Ver commits

# Carregar LKM para testes
cd DEVBRAIN_V23/kernel/lkm
sudo insmod devbrain_ai.ko
lsmod | grep devbrain_ai

# Rodar tests
pytest tests/test_kernel_ai.py -v

# Ver documentação
cat docs/KERNEL_INTEGRATION.md
```

PRÓXIMAS FASES (After Phase 10)

Após completar Phase 10, você estará pronto para:

- **Phase 11:** Consciência Aprofundada + Psicoanálise
- **Phase 12:** Produção Hardened (24/7 uptime)
- **Phase 13:** Evolução Contínua

STATUS FINAL: READY FOR EXECUTION ✓

Comece AGORA. Seu Copilot está esperando.